

New lms frontend and staged computation graphs

Ruben Fiszel

December 7, 2016

~ Supervised by Nada Amin and Prof. Martin Odersky ~



ÉCOLE POLYTECHNIQUE
FÉDÉRALE DE LAUSANNE

Abstract

In this report, we explore staging, in particular the LMS framework and the development of its new frontend whose aim is ease the writing of staged dsl through, among others, shadowing of types. We also explore the usage of this new frontend for a particular case study: Staged computation graphs.

Contents

Introduction	3
1 LMS	5
Staging	8
Why staging	8
Exp tree	8
Deep reuse of embedding language order	8
Frontend/Backend ?	8
lms	8
library author	8
delite	8
user	8
2 The new frontend	9
Lift	9
Typeclass	9
Typeclass overloading	9
Primitives types and collections	9
3 Computation Graph	10
Graph	11
Cycle check	11
Arithmetic	11
Benchmark	11
DerivableGraph	11

Backpropagation	11
MatrixGraph	11
Dimensions check	11
Conclusion	12
Acknowledgement	13

Introduction

If programming can be considered as an art, then programming languages are the brush of the artists. The tools are chosen according to individual preferences, and of course the constraints of the desired final work. One of those major constraint is efficiency. Efficiency is less crucial nowadays than it was back when computing power was expensive and scarce, but it is still a very desirable goal. Efficiency can be achieved by writing explicitly very precisely each step of the program. To enable this, some programming languages are relatively close from the model of the hardware. Those programming languages are called “low-level”, because close from the machine and thus distant from the heights of abstraction. But as programs grow more complex, the need for abstraction does too.

To raise abstraction power, expressive languages with complex compilers were made. Ironically, the first programming language, lambda-calculus, was the epitome of abstraction and thought of before the first very computer to execute it. Those languages, mostly from the functional programming world, aim the user to express the intent than the steps, the what rather than the how. The compilers needed for them are heavy machinery that applies all sort of optimization before translating source code to machine language. The more abstract the language, the more gap there is to fill, and the more compilers have opportunity to optimize.

Nevertheless, growing complexity in compiler is not the panacea. In a wide range of programs, compilers are limited by the lack of domain specific knowledge of the program. Indeed, constraining a program to one specific domain open the door for a wide range of specific optimization. Furthermore, code generation by “macro” is a rather poor way to enable the further abstraction brought by code that generate code.

One solution to both issue could be to extend the compiler for each domain and having a very powerful macro system. The other one is to write specific DSL in a staged environment. It is the latter that is explored with LMS, a scala library for runtime code generation. In this report, we will explore a new lms implementation that brings a user frontend, more convenient for the end-user and with multiple benefits enabled by extended typeclass usage. We also cover a case study of lms applied to the domain of computation graphs.

1. LMS

Lightweight Modular Staging (LMS) is a library written in Scala that enable staged meta-programming. Meta-Programming is the art of writing computer programs that can handle other programs as data. Thus, they are able to reason about programs, which they are themselves, hence the name. In staged meta-programming, the goal is to generate new programs from annotated programs. The generation can go through multiple stage. Each stage is another compilation time that can leverage new informations to build the final program.

One clear example of this is sparse vector multiplication.

```
val v1: IndexedSeq
val block = {
  val prng = new util.Random(123)
  def rd(i: Int) = prng.nextInt(i)

  def newVec(size: Int) = {
    val a = Array.fill(size)(0)
    for (i <- (1 to 10))
      a(rd(size)) = rd(100)
```

```

      a
    }

    val v2 = newVec(1000)

    v1.zip(v2) map { case (x1, x2) => x1 * x2 } sum
  }

```

This work but it is terribly inefficient. What about all those 0 multiplications. Surely, we can do something better. There is 3 cases.

- v1 and v2 are both known at compile time
- v1 is known at compile time
- v1 and v2 are both only known at runtime

In all three cases, LMS can generate efficient code that avoid the (1000 [zip]+ 1000 [*]) operations at runtime and reduce it to 10 operations (no more zip) and even 0 operations in the first case!

How ? Well it is all about available information. We will start with the first case which is the simplest to understand: We know at staging time all the data and all the operation that is applied to the data. Instead of waiting for runtime to apply those operations, we can generate a new block that is strictly equivalent. *Note* that to ensure equivalence, we use a fixed seed so that the number generation is deterministic and the current state of the world at execution is not relevant. So the generated code can be reduced to:

```

val stagedBlock = C

```


where C is a constant which value depend on $v1$

In the second and third case the code can be generated to this reduced form:

```
//TODO
```

```
val stagedBlock = {  
    val x0 = new util.Random(123)  
    val x1 = v1(x0.nextInt(100)) * rd(100)  
    val x2 = v1(rd(100)) * rd(100)  
    val x3 = x2 + x1  
    val x4 = v1(rd(100)) * rd(100)  
}
```

To be continued

Staging

Why staging

Exp tree

Deep reuse of embedding language order

Frontend/Backend ?

lms

library author

delite

user

2. The new frontend

Lift

Typeclass

Typeclass overloading

Primitives types and collections

3. Computation Graph

Graph

Cycle check

Arithmetic

Benchmark

DerivableGraph

Backpropagation

MatrixGraph

Conclusion

blablabla

Acknowledgement

Thanks to my beloved parents, my awesome supervisor Nada Amin, Prof. Martin Odersky, the lms master and author Tiark Rumpf, and the delite folks Kevin James Brown and David Koeplinger.